

## Question 1. Pseudocode for Multi-Level Queue Dispatcher

1. Initialise the queues (RT job dispatch Q, Normal job dispatch Q, **Level 0 Q**, **Level 1 Q** and **Level 2 Q**)
2. Populate RT job dispatch Q and Normal job dispatch Q from the input file
3. Ask the user to enter an integer value for **time\_quantum**
4. While there is a currently running process or any of the queues are not empty
  - i. Unload all arrived processes into their respective Level Qs
  - ii. If there is a process running
    - a. If the priority value is **0** (**Level 0** process running)
      - A. **SIGINT** to terminate process
      - B. Calculate turnaround\_time and wait\_time for process
      - C. Free process struct memory
    - b. If the priority value is **1** (**Level 1** process running)
      - A. Decrease <remaining\_cpu\_time> by **time\_quantum**
      - B. If time has been exhausted
        - **SIGINT** to terminate process
        - Calculate turnaround\_time and wait\_time for process
        - Free process struct memory
      - C. Else
        - **SIGTSTP** to suspend process
        - Set priority to **2**
        - Enqueue to tail of **Level 2 Q**
    - c. If the priority value is **2** (**Level 2** process running)
      - A. Decrease <remaining\_cpu\_time> by 1
      - B. If the time has been exhausted
      - C. If there is another process in **Level 0 Q**
        - **SIGTSTP** to suspend process
        - Enqueue to head of **Level 2 Q**
        - Dequeue process from the head of **Level 0 Q**
        - Start and set as the currently running process
      - D. If there is another process in **Level 1 Q**
        - **SIGTSTP** to suspend process
        - Enqueue to head of **Level 2 Q**
        - Dequeue process from the head of **Level 1 Q**
        - Start and set as the currently running process
      - E. Else
        - If time has been exhausted
          - **SIGINT** to terminate process
          - Calculate turnaround\_time and wait\_time for process
          - Free process struct memory

- iii. If there is no process running and at least one of the Level Qs aren't empty (we schedule the next job with highest priority)
  - a. If **Level 0 Q** is not empty
    - A. Dequeue process from the head of **Level 0 Q**
    - B. Start and set as the currently running process
  - b. Else if **Level 1 Q** is not empty
    - A. Dequeue process from the head of **Level 1 Q**
    - B. Start and set as the currently running process
  - c. Else if **Level 2 Q** is not empty
    - A. Dequeue process from the head of **Level 2 Q**
    - B. If the process is a suspended process
      - Send **SIGCONT** to resume it
      - and set as the currently running process
    - C. Else
      - Start and set as the currently running process
- iv. If there is a current process
  - a. Sleep for **decrease\_time** (different for each level process)
  - b. Increase timer by **decrease\_time**
- v. Else
  - a. Sleep for 1
  - b. Increase timer by 1
- vi. Go back to 4.

## Question 2. Example and Explanation of 3 'Job Dispatch List Files' and their Outputs

### Example 1. Testing **Level 0 Q**

The purpose of this example is to **test that the Level 0 Q is always prioritised**. For this example we assume a **time\_quantum** of 3 and then compare the results with 4 and 5.

Process	<arrival time>	<cpu_time>	<priority>
1	0	5	1
2	4	2	0
3	6	3	1
4	9	4	0
5	10	3	1
6	11	2	1

When this example is first read in, assuming that the initial initialisation of queues are correct, process 1 will already have been dispatched and enqueued onto the **Level 1 Q** since it has an <arrival time> of 0. It will then be run for 3 seconds (**time\_quantum**), then pre-empted and enqueue onto the **Level 2 Q**. Since process 1 is at the front of the Level 2 Q and there are no other processes with higher priority, it will run for 1 second until process 2 arrives. Here the program should preempt process 1 and allow process 2 to run until completion as its priority is 0. Then process 1 waits until all the other processes finish running as they all have higher priority, and none of them ever get appended to the level 2 Q because they all have a <cpu\_time> of less than **time\_quantum**. Finally process 1, the last process, runs for the <remaining\_cpu\_time> of 1 second.

A **time\_quantum** of 3 is selected to ensure that process 2, the job with higher priority, interrupts and enforces the program to preempt process 1 and start running process 2 at the time of its arrival. Furthermore, it ensures that process 4 runs until completion without interruption even though the **time\_quantum** is less than its <cpu\_time>, ensuring that all jobs with priority 0 are always prioritised.

When running this job list with a **time\_quantum** of 5 or greater, the order of jobs remains in a FCFS fashion hence, no preempting is evident thus, we cannot test for priority in level 0.

#### **Expected outcomes:**

- **Time Quantum: 3 & 4 (both give the same results)**
  - Order of finishing jobs → 2,3,4,5,6,1
  - Average Turnaround Time → 6.83 seconds
  - Average Wait Time → 3.67 seconds
- **Time Quantum: 5**
  - Order of finishing jobs → 1,2,3,4,5,6
  - Average Turnaround Time → 5.33 seconds
  - Average Wait Time → 2.17 seconds

## Example 2. Testing Level 2 Q

The purpose of this example is to **test that the Level 2 Q enqueues jobs back on the head of the queue rather than the tail**. For this example we assume a **time\_quantum** of 3 and then compare the results with 4.

Process	<arrival time>	<cpu_time>	<priority>
1	0	5	1
2	1	5	1
3	6	1	0
4	7	1	0
5	8	1	0
6	10	3	1

When this example is first read in, assuming that the initial initialisation of queues are correct, process 1 will already have been dispatched and enqueued onto the **Level 1 Q** since it has an <arrival time> of 0. It will then be run for 3 seconds (**time\_quantum**), then pre-empted and enqueue onto the **Level 2 Q**. Next, just like process 1, process 2 runs for 3 seconds and gets pre-empted onto the **Level 2 Q**. Furthermore, processes 3, 4 and 5 will run to completion in that order because they have priority 0. Upon completion, process 6 hasn't arrived yet however process 1 is at the head of the **Level 2 Q**, therefore, it can run for 1 second before process 6 arrives. After running for 1 second it should enqueue it back onto the head of the **Level 2 Q** (**this is where the test for this specific example occurs**). After process 6 finishes running, process 1 should run again to completion and finally, process 2 should then finish running.

A **time\_quantum** of 3 is selected to ensure the test of whether the process gets enqueued back onto the head of the Level 2 can be tested. Any **time\_quantum** larger or smaller than 3 will not be able to test this property.

### Expected outcomes:

- **Time Quantum: 3**
  - Order of finishing jobs → 3,4,5,6,1,2
  - Average Turnaround Time → 5.83 seconds
  - Average Wait Time → 3.17 seconds
- **Time Quantum: 4**
  - Order of finishing jobs → 3,4,5,6,1,2
  - Average Turnaround Time → 7.17 seconds
  - Average Wait Time → 4.50 seconds

Although the order of finishing jobs are the same, the property is clearly not tested.