



*Operating Systems:
Internals and Design Principles*
William Stallings

Chapter 4
Threads





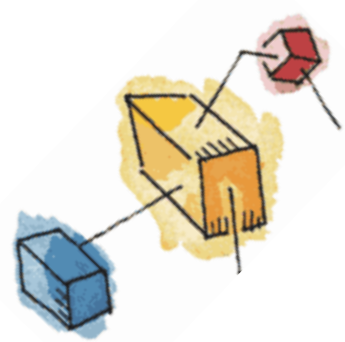
Processes and Threads

- Processes have two characteristics:
 - **Resource ownership** - process includes a virtual address space to hold the process image
 - the OS performs a protection function to prevent unwanted interference between processes with respect to resources
 - **Scheduling/execution** - follows an execution path that may be interleaved with other processes
 - a process has an execution state (Running, Ready, etc.) and a dispatching priority and is scheduled and dispatched by the OS
- These two characteristics are treated independently by modern operating systems:
 - The unit of dispatching is referred to as a **thread** or lightweight process
 - The unit of resource ownership is referred to as a **process** or **task**



Process


- The unit of resource allocation and a unit of protection
- A process is associated with
 - A virtual address space which holds the process image
 - Protected access to
 - Processors,
 - Other processes,
 - Files,
 - I/O resources





Multiple Threads in Process

- Each thread has
 - Access to the memory and resources of its process (all threads of a process share this)
 - An execution state (running, ready, etc.)
 - Saved thread context when not running
 - An execution stack
 - Some per-thread static storage for local variables





Threads vs. processes

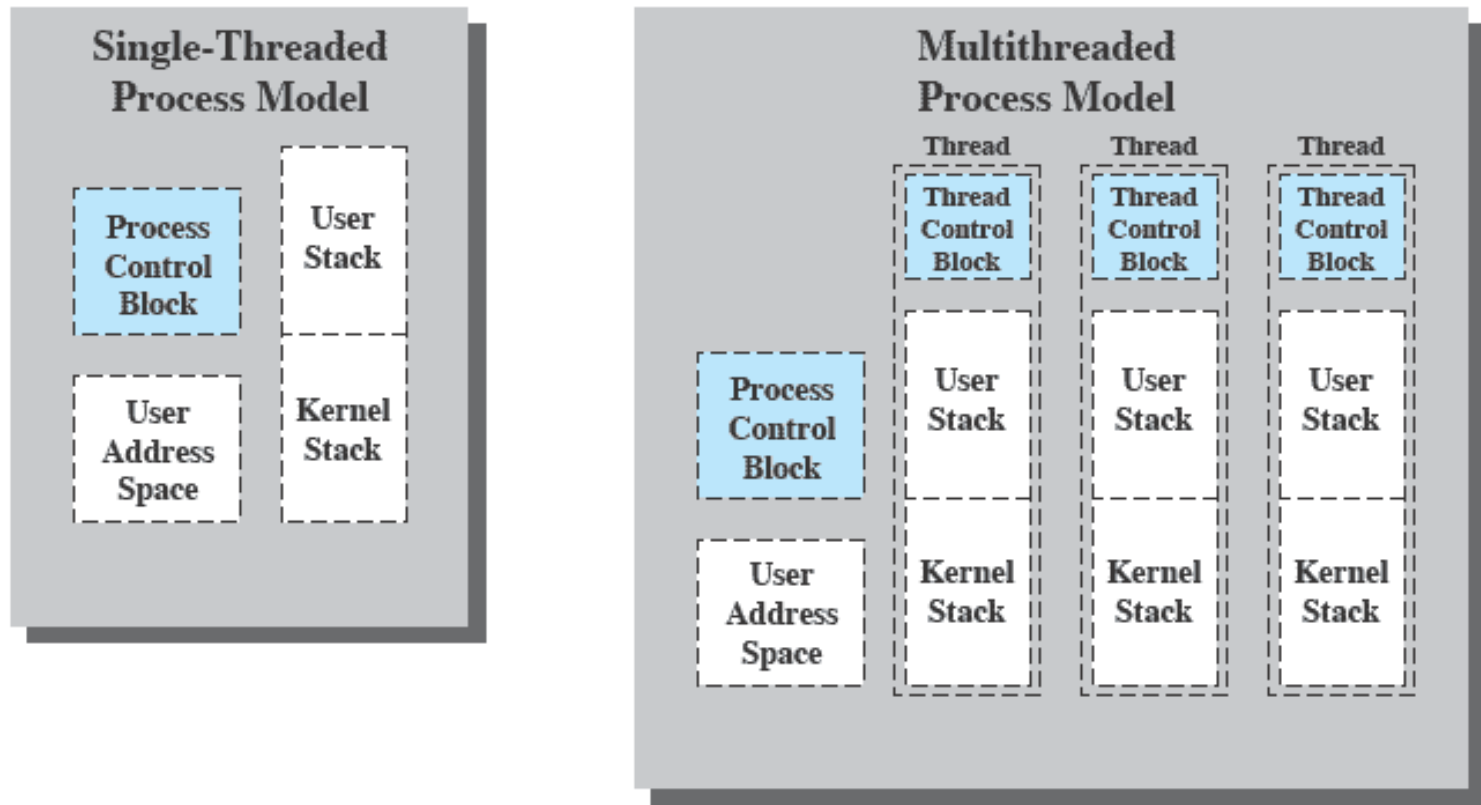
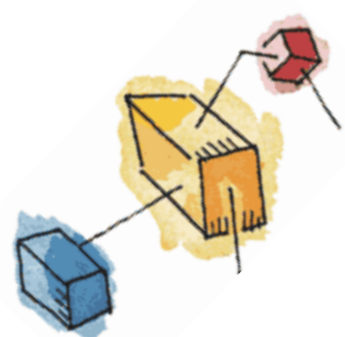


Figure 4.2 Single Threaded and Multithreaded Process Models






Example: Remote Procedure Call

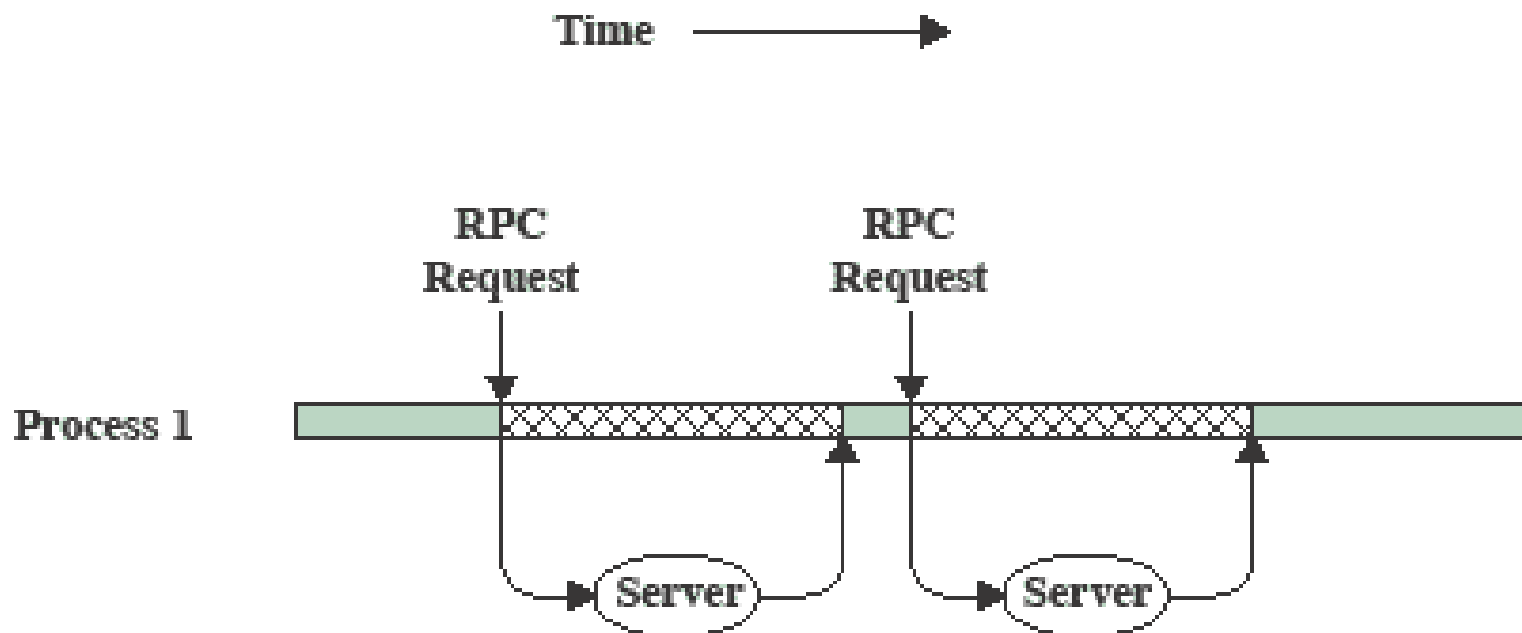
- Consider:
 - A program that performs two remote procedure calls (RPCs)
 - to two different hosts
 - to obtain a combined result.





RPC


Using Single Thread






(a) RPC Using Single Thread




RPC Using One Thread per Server



(b) RPC Using One Thread per Server (on a uniprocessor)


-  **Blocked, waiting for response to RPC**
-  **Blocked, waiting for processor, which is in use by Thread B**
-  **Running**



Thread Execution States

- Three basic states: Running, Ready, and Blocked
- Operations associated with a change in thread state
 - Spawn (another thread)
 - Allocate register context and stacks
 - Block
 - moved to an event queue waiting for the event
 - Issue: Will blocking a thread block other, or *all*, threads within the same process?
 - Unblock
 - moved to the Ready queue for execution
 - Finish (thread)
 - De-allocate register context and stacks





Thread Synchronization

- It is necessary to synchronize the activities of the various threads so that they do not interfere with each other
 - all threads of a process share the same address space and other resources
 - any alteration of a resource by one thread affects the other threads in the same process
- In general, the techniques used for thread synchronization are the same as those for process synchronization





Types of Threads



- User Level Thread (ULT)
- Kernel level Thread (KLT) also called:
 - kernel-supported threads
 - lightweight processes.






User-Level Threads

- All thread management is done by the application
 - Multithreading is managed by a runtime threads library
- The kernel is not aware of the existence of threads
 - Only schedule the process as a unit and assigns a single execution state to that process






Disadvantages of ULT


- In a typical OS many system calls are blocking
 - as a result, when a ULT executes a system call, not only is that thread blocked, but all of the threads within the process are blocked
- In a pure ULT strategy, a multithreaded application cannot take advantage of multiprocessing





Kernel-Level Threads

- Kernel maintains context information for the process and the threads
 - No thread management done by application
- Scheduling is done on a thread basis



(b) Pure kernel-level





Kernel-Level Threads

- In an OS that supports threads, scheduling and dispatching is done on a thread basis
- Most of the state information dealing with execution is maintained in thread-level data structures
- However, several actions will affect all of the threads in a process and the OS must manage at the process level
 - Suspending a process involves suspending all threads of the process
 - Termination of a process, terminates all threads within the process






Advantages of KLT


- The kernel can simultaneously schedule multiple threads from the same process on multiple processors.
- If one thread in a process is blocked, the kernel can schedule another thread of the same process.





Multicore & Multithreading

single core



core 1




core 2



time






Multicore & Multithreading

- Achieves concurrency without the overhead of using multiple processes
 - Threads within the same process can exchange information through their common address space and have access to the shared resources of the process





Multicore & Multithreading


- Multicore systems putting pressure on programmers, challenges include
 - Dividing activities
 - Balance
 - Data dependency
 - Testing and debugging



Linux Threads

- Linux uses the same internal representation for processes and threads; a thread is simply a new process (or task) that happens to share the same address space as its parent
- A distinction is only made when a new thread is created by the `clone` system call
 - `fork` creates a new process with its own entirely new process context
 - `clone` creates a new process with its own identity, but that is allowed to share the data structures of its parent



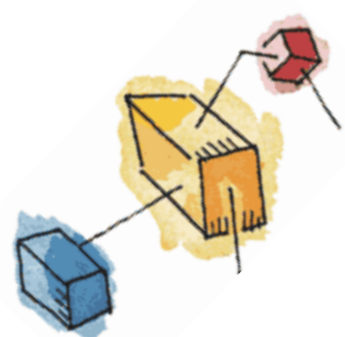


Linux Threads

- Using `clone` gives an application fine-grained control over exactly what is shared between two threads

flag	meaning
<code>CLONE_FS</code>	File-system information is shared.
<code>CLONE_VM</code>	The same memory space is shared.
<code>CLONE_SIGHAND</code>	Signal handlers are shared.
<code>CLONE_FILES</code>	The set of open files is shared.





Windows Threads

- Windows makes use of two types of process-related objects:
- Processes
 - an entity corresponding to a user job or application that owns resources
- Threads
 - a dispatchable unit of work that executes sequentially and is interruptible





Process & Thread Objects

Object Type

Process

Object Body
Attributes

Process ID
Security Descriptor
Base priority
Default processor affinity
Quota limits
Execution time
I/O counters
VM operation counters
Exception/debugging ports
Exit status

Services

Create process
Open process
Query process information
Set process information
Current process
Terminate process

(a) Process object

Object Type

Thread

Object Body
Attributes

Thread ID
Thread context
Dynamic priority
Base priority
Thread processor affinity
Thread execution time
Alert status
Suspension count
Impersonation token
Termination port
Thread exit status

Services

Create thread
Open thread
Query thread information
Set thread information
Current thread
Terminate thread
Get context
Set context
Suspend
Resume
Alert thread
Test thread alert
Register termination port

(b) Thread object





Summary

- Process/related to resource ownership
- Thread/related to program execution
- User-level threads
 - created and managed by a threads library that runs in the user space of a process
 - a mode switch is not required to switch from one thread to another
 - only a single user-level thread within a process can execute at a time
 - if one thread blocks, the entire process is blocked
- Kernel-level threads
 - threads within a process that are maintained by the kernel
 - a mode switch is required to switch from one thread to another
 - multiple threads within the same process can execute in parallel on a multiprocessor
 - blocking of a thread does not block the entire process

