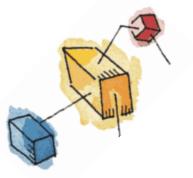
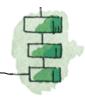
Operating Systems: Internals and Design Principles William Stallings

Chapter 5
Concurrency: Mutual
Exclusion and
Synchronization (cont.)



#### Outline

- OS synchronization mechanisms
  - Semaphore
    - Binary and general semaphores
    - Implementation
  - Producer/consumer problem
  - Bounded buffer
  - Monitor
    - Condition variables
  - Message passing
  - Reader/Writer problem



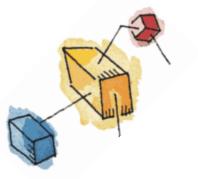


## Common Concurrency Mechanisms

Semaphore	An integer value used for signaling among processes. Only three operations may be performed on a semaphore, all of which are atomic: initialize, decrement, and increment. The decrement operation may result in the blocking of a process, and the increment operation may result in the unblocking of a process. Also known as a counting semaphore or a general semaphore	
Binary Semaphore	A semaphore that takes on only the values 0 and 1.	
Mutex	Similar to a binary semaphore. A key difference between the two is that the process that locks the mutex (sets the value to zero) must be the one to unlock it (sets the value to 1).	
Condition Variable	A data type that is used to block a process or thread until a particular condition is true.	
Monitor	A programming language construct that encapsulates variables, access procedures and initialization code within an abstract data type. The monitor's variable may only be accessed via its access procedures and only one process may be actively accessing the monitor at any one time. The access procedures are <i>critical sections</i> . A monitor may have a queue of processes that are waiting to access it.	
Event Flags	A memory word used as a synchronization mechanism. Application code may associate a different event with each bit in a flag. A thread can wait for either a single event or a combination of events by checking one or multiple bits in the corresponding flag. The thread is blocked until all of the required bits are set (AND) or until at least one of the bits is set (OR).	
Mailboxes/Messages	A means for two processes to exchange information and that may be used for synchronization.	
Spinlocks	Mutual exclusion mechanism in which a process executes in an infinite loop waiting for the value of a lock variable to indicate availability	



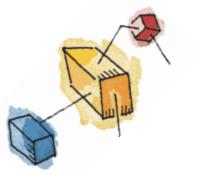




#### Semaphore

- A queue is used to hold processes waiting on the semaphore – eliminating busy-wait
- Semaphore:
  - An integer value used for signalling among processes.
- Only three operations may be performed on a semaphore, all of which are atomic:
  - Initialize the integer,
  - decrement the value (semWait),
  - increment the value (semSignal)





#### Semaphore

- The semaphore s is initially assigned a zero or positive value
- When a process issues a semWait, s is decremented
  - The process will be blocked if s goes negative
  - The negative value equals the number of blocked processes waiting to be unblocked
- Each semSignal, incrementing s, unblocks one of the waiting processes when s is negative



#### Semaphore Primitives

```
struct semaphore {
     int count;
     queueType queue;
void semWait(semaphore s)
     s.count--;
     if (s.count < 0) {
          /* place this process in s.queue */;
          /* block this process */;
void semSignal(semaphore s)
     s.count++;
     if (s.count <= 0) {
          /* remove a process P from s.queue */;
          /* place process P on ready list */;
```



Figure 5.3 A Definition of Semaphore Primitives



## **Binary Semaphore Primitives**

```
struct binary semaphore {
     enum {zero, one} value;
     queueType queue;
void semWaitB(binary semaphore s)
     if (s.value == one)
          s.value = zero;
     else {
             /* place this process in s.queue */;
             /* block this process */;
void semSignalB(semaphore s)
     if (s.queue is empty())
          s.value = one;
     else {
             /* remove a process P from s.queue */;
             /* place process P on ready list */;
```



Figure 5.4 A Definition of Binary Semaphore Primitives

## plementation of Semaphores

- the semWait and semSignal operations themselves are critical sections and thus must be implemented as atomic primitives
- Use one of the hardware-supported schemes for mutual exclusion





## Strong/Weak Semaphores

- A queue is used to hold processes waiting on the semaphore – eliminating busy-wait
  - In what order are processes removed from the queue?

#### Strong Semaphores

• the process that has been blocked the longest is released from the queue first (FIFO)

#### Weak Semaphores

• the order in which processes are removed from the queue is not specified



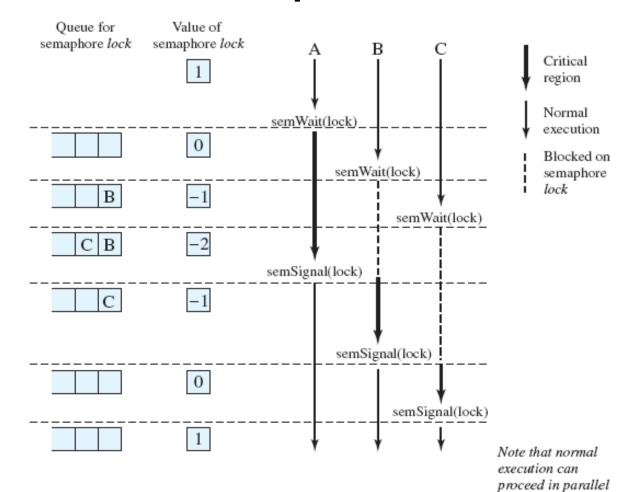
## Mutual Exclusion Using Semaphores

```
/* program mutualexclusion */
const int n = /* number of processes */;
semaphore s = 1; /* Initialize s to 1 */
void P(int i)
     while (true) {
          semWait(s);
          /* critical section */;
          semSignal(s);
          /* remainder */;
void main()
     parbegin (P(1), P(2), . . ., P(n));
```



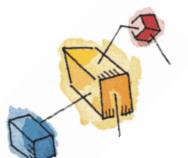


# Mutual Exclusion Using Semaphores





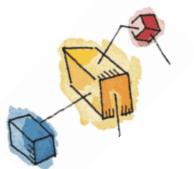
but that critical regions are serialized.



# Synchronization Using Semaphores

```
Semaphore s = 0; /* Initialize s to 0 */
Proc 0() {
                         proc 1() {
                          semWait (s);
 S1;
 semSignal (s);
                          S2;
```





# Producer/Consumer Problem

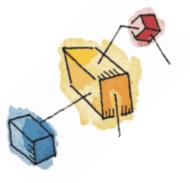
#### General Situation:

- One or more producers are generating data and placing these in a buffer
- One or more consumers are taking items out of the buffer one at time
- Only one producer or consumer may access the buffer at any one time

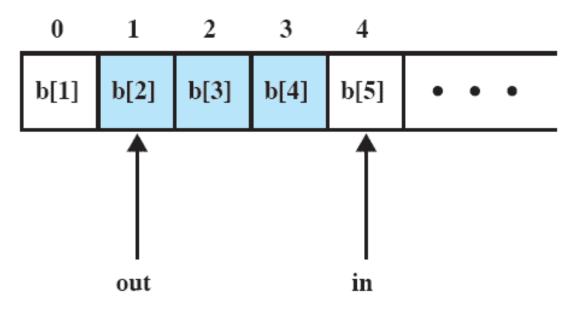
#### • The Problem:

 Ensure that the Producer can't add data into full buffer and consumer can't remove data from empty buffer





#### **Buffer Structure**

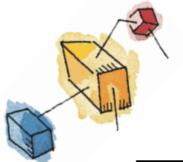


Note: shaded area indicates portion of buffer that is occupied

Figure 5.8 Infinite Buffer for the Producer/Consumer Problem







### **Binary Semaphores: Incorrect Solution**

```
/* program producerconsumer */
                 int n;
                binary semaphore s = 1, delay = 0;
                void producer()
                      while (true) {
                           produce();
                           semWaitB(s);
                           append();
                           n++;
                           if (n==1) semSignalB(delay);
                           semSignalB(s);
                void consumer()
                      semWaitB(delay);
                      while (true) {
                           semWaitB(s);
                           take();
A producer can update n
                           n--;
                           semSignalB(s);
before if statement!
                          consume();
                           if (n==0) semWaitB(delay);
                void main()
                      parbegin (producer, consumer);
```





# Binary Semaphores: Correct Solution

```
/* program producerconsumer */
         int n:
        binary semaphore s = 1, delay = 0;
        void producer()
                                       Can cause starvation!
              while (true) {
                   produce();
                   semWaitB(s);
Producers continue
                   append();
                   n++;
to append &
                    if (n==1) semSignalB(delay);
make n > 1 always!
                   semSignalB(s);
        void consumer()
C1
              int m; /* a local variable */
            semWaitB(delay);
delav = 0
              while (true) {
Starved here!
                   semWaitB(s);
                   take();
                   n--;
                   m = n;
                   semSignalB(s);
                    consume();
        C0:
                  if (m==0) semWaitB(delay);
        m > 0
        void main()
              n = 0;
              parbegin (producer, consumer);
```





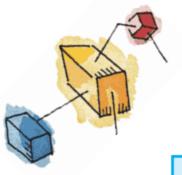
#### General Semaphores

buffer is empty

```
nsure consumer not /* program producerconsumer */
to take items when the semaphore n = 0, s = 1; void producer()
                                                           For mutual exclusion
                                while (true) {
                                      produce();
                                      semWait(s);
                                      append();
                                      semSignal(s);
                                      semSignal(n);
                           void consumer()
                                while (true) {
                                      semWait(n);
                                      semWait(s);
                                      take();
                                      semSignal(s);
                                      consume();
                           void main()
                                 parbegin (producer, consumer);
```

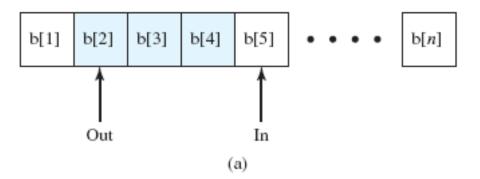


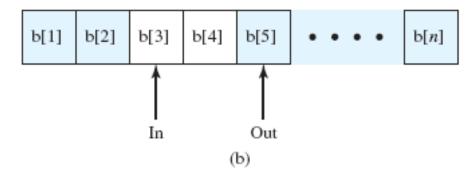
Figure 5.11 A Solution to the Infinite-Buffer Producer/Consumer Problem Using Semaphores

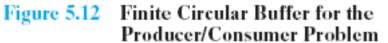


#### **Bounded Buffer**

Block on:	Unblock on:
Producer: insert in full buffer	Consumer: item inserted
Consumer: remove from empty buffer	Producer: item removed

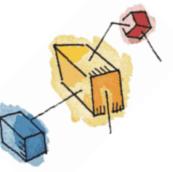












#### Semaphores

```
/* program boundedbuffer */
                  const int sizeofbuffer = /* buffer size */;
                  semaphore s = 1, n= 0, e= sizeofbuffer;
                  void producer()
                                                           Ensure producer not
                       while (true) {
                                                           to add items when the
                             produce();
 Producers wait here if
                                                           buffer is full
                             semWait(e);
 the buffer is full
                             semWait(s);
                             append();
Producer informs
                             semSignal(s);
Consumers that there
                             semSignal(n);
are data items available
                  void consumer()
                       while (true) {
                             semWait(n);
                             semWait(s);
                             take();
                             semSignal(s);
                             semSignal(e);
                             consume();
                  void main()
                       parbegin (producer, consumer);
```



## **İssues with Semaphores**

- Semaphores provide a powerful synchronization tool. However,
  - semSignal() and semWait() are scattered among several processes. Therefore, it is difficult to understand their effects
  - Usage must be correct in all the processes.
  - One bad process (or one programming error) can kill the whole system.





## **İssues with Semaphores**

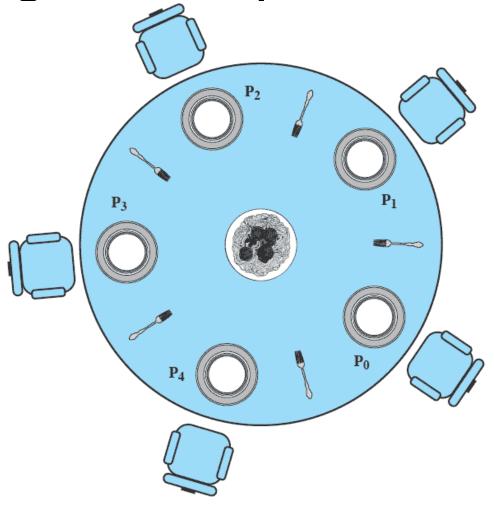
- Deadlock two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

```
P_0 P_1 semWait (S); semWait (Q); semWait (Q); semWait (Q); ... ... semSignal (S); semSignal (Q); semSignal (Q);
```

 Starvation – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.



pining Philosophers Problem



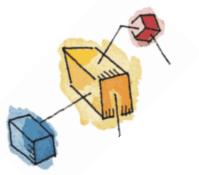


### pining Philosophers Problem

```
diningphilosophers */
/* program
semaphore fork [5] = {1};
int i;
void philosopher (int i)
     while (true) {
                                     Warning: This solution could
          think();
                                     create a deadlock!
          wait (fork[i]);
          wait (fork [(i+1) mod 5]);
          eat();
          signal(fork [(i+1) mod 5]);
          signal(fork[i]);
void main()
    parbegin (philosopher (0), philosopher (1), philosopher
(2)_{r}
          philosopher (3), philosopher (4));
```



Figure 6.12 A First Solution to the Dining Philosophers Problem



#### **Monitors**

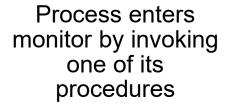
- The monitor is a programming-language construct that provides equivalent functionality to that of semaphores and that is easier to control.
- Implemented in a number of programming languages, including
  - Concurrent Pascal, Pascal-Plus, Modula-2, Modula-3, and Java
- Software module consisting of one or more procedures, an initialization sequence, and local data



#### **Monitor Characteristics**

Local data variables are accessible only by the monitor's procedures and not by any external procedure

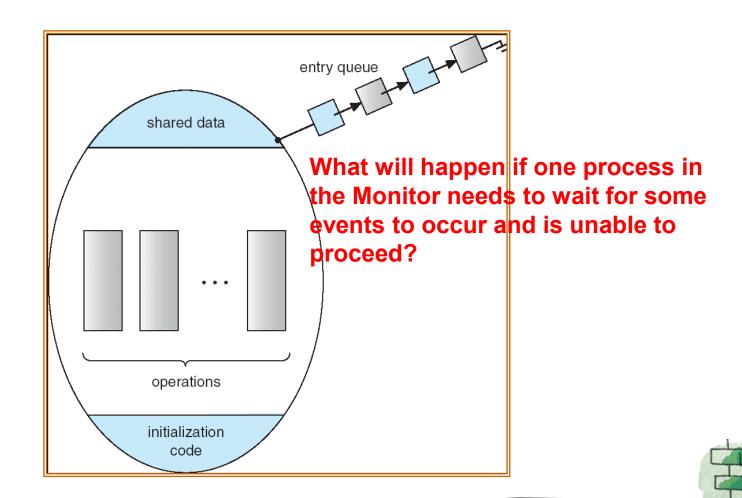
Only one process may be executing in the monitor at a time



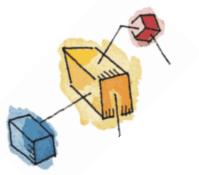




## Schematic View of a Monitor





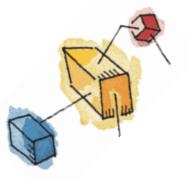


### Synchronization

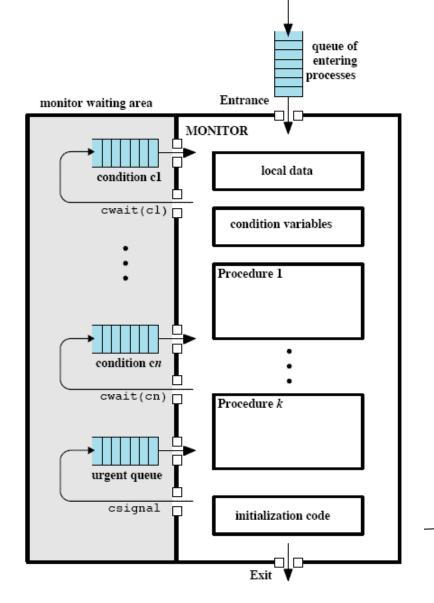
- Synchronisation achieved by using condition variables that are contained within a monitor and only accessible within the monitor
- Condition variables are operated on by two functions:
  - cwait(c): Suspend execution of the calling process on condition c
  - csignal(c) Resume execution of some process blocked after a cwait(c) on the same condition







#### Structure of a Monitor





# Bounded Buffer Solution Using Monitor

```
/* program producerconsumer */
monitor boundedbuffer;
char buffer [N];
                                                     /* space for N items */
int nextin, nextout;
                                                      /* buffer pointers */
int count;
                                             /* number of items in buffer */
cond notfull, notempty; /* condition variables for synchronization */
void append (char x)
    if (count == N) cwait(notfull);
                                     /* buffer is full; avoid overflow */
    buffer[nextin] = x;
    nextin = (nextin + 1) % N;
    count++;
    /* one more item in buffer */
                                          /* resume any waiting consumer */
    csignal(notempty);
void take (char x)
    if (count == 0) cwait(notempty); /* buffer is empty; avoid underflow */
    x = buffer[nextout];
    nextout = (nextout + 1) % N;
                                              /* one fewer item in buffer */
    count--;
                                           /* resume any waiting producer */
    csignal(notfull);
                                                          /* monitor body */
    nextin = 0; nextout = 0; count = 0; /* buffer initially empty */
```

```
void producer()
    char x;
    while (true) {
    produce(x);
    append(x);
void consumer()
    char x;
    while (true) {
      take(x);
      consume(x);
void main()
    parbegin (producer, consumer);
```



#### Message Passing

When processes interact with one another, there are two fundamental requirements:

#### synchronization

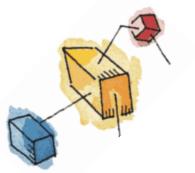
to enforce mutual exclusion

#### communication

to exchange information

- Message Passing is one approach to providing these functions
  - works with distributed systems and shared memory multiprocessor and uniprocessor systems





### Message Passing

- The actual function of message passing is normally provided in the form of a pair of primitives:
  - send (destination, message)
  - receive (source, message)
- Exchange information
  - A process sends information in the form of a message to another process designated by a destination
  - A process receives information by executing the receive primitive, indicating the source and the message





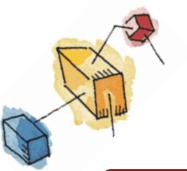


## Blocking send, Blocking receive

- Both sender and receiver are blocked until message is delivered
- Synchronization
  - Receiver cannot proceed until the message is received
  - Sender cannot proceed until the message arrives to the destination







## Non-blocking Send, Non-blocking Receive

#### Nonblocking send, blocking receive

- sender continues on but receiver is blocked until the requested message arrives
- most useful combination
- sends one or more messages to a variety of destinations as quickly as possible
- example -- a service process that exists to provide a service or resource to other processes

#### Nonblocking send, nonblocking receive

• neither party is required to wait





#### Addressing

- Sending process need to be able to specify which process should receive the message
  - Direct addressing
  - Indirect Addressing







#### **Direct Addressing**

- Send primitive includes a specific identifier of the destination process
- Receive primitive can be handled in one of two ways:
  - require that the process explicitly designate a sending process
    - effective for cooperating concurrent processes
  - implicit addressing
    - source parameter of the receive primitive possesses a value returned when the receive operation has been performed





### Indirect addressing

Messages are sent to a shared data structure consisting of queues that can temporarily hold messages



Queues are referred to as *mailboxes* 



Allows for greater flexibility in the use of messages



One process sends a message to the mailbox and the other process picks up the message from the mailbox







# Indirect Process Communication

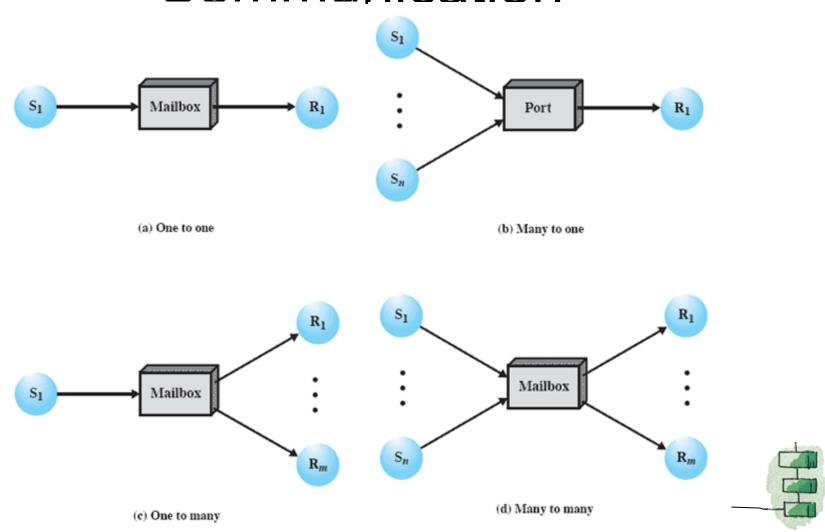
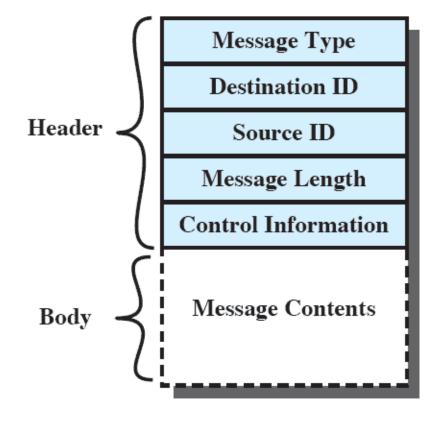
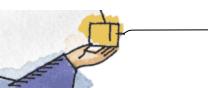


Figure 5.18 Indirect Process Communication

## General Message Format

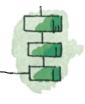




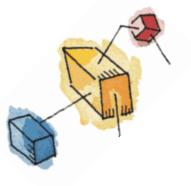


## Readers/Writers Problem

- A data area is shared among many processes
  - Some processes only read the data area (readers),
     some only write to the area (writers)
- Conditions that must be satisfied:
  - Multiple readers may read the file at once.
  - 2. Only one writer at a time may write
  - 3. If a writer is writing to the file, no reader may read it.







#### Readers have Priority

Readers continue to arrive & make readcount > 1 always!

Writers may get starved here!

```
/* program readersandwriters */
int readcount;
                                    To ensure
semaphore x = 1, wsem = 1;
                                    1. only one writer can write at
void reader()
                                       a time
   while (true) {
                                    2. no writer can write when
     semWait (x);
                                       there are readers
     readcount++;
     if (readcount == 1) semWait (wsem);
     semSignal (x);
     READUNIT();
     semWait (x);
     readcount --;
     if (readcount == 0) semSignal (wsem);
     semSignal (x);
void writer()
   while (true) {
     semWait (wsem);
     WRITEUNIT();
     semSignal (wsem);
void main()
    readcount = 0;
    parbegin (reader, writer);
```







### Writers have Priority

```
/* program readersandwriters */
int readcount, writecount;
semaphore x = 1, y = 1, z = 1, wsem = 1, rsem = 1;
void reader()
                                    Readers and writers
   while (true) {
                                    compete for "rsem"
    semWait (z);
         semWait (rsem);
              semWait (x);
                   readcount++;
                   if (readcount == 1) semWait (wsem);
              semSignal (x);
         semSignal (rsem);
    semSignal (z);
    READUNIT();
    semWait (x);
         readcount--;
         if (readcount == 0) semSignal (wsem);
    semSignal (x);
```

```
void writer ()
   while (true) {
     semWait (y);
          writecount++;
    → if (writecount == 1) semWait (rsem);
     semSignal (y);
     semWait (wsem);
     WRITEUNIT();
     semSignal (wsem);
     semWait (y);
          writecount--;
          if (writecount == 0) semSignal (rsem);
     semSignal (y);
void main()
   readcount = writecount = 0;
   parbegin (reader, writer);
```



## Writers have Priority

```
/* program readersandwriters */
int readcount, writecount;
semaphore x = 1, y = 1, z = 1, wsem = 1, rsem = 1;
void reader()
                     "z" is used to make sure
   while (true) { only one reader is competing
    semWait (z) "rsem" with writers
         semWait (rsem);
             semWait (x);
                  readcount++;
                  if (readcount == 1) semWait (wsem);
             semSignal (x);
         semSignal (rsem);
    semSignal (z);
    READUNIT();
    semWait (x);
         readcount--;
         if (readcount == 0) semSignal (wsem);
    semSignal (x);
```

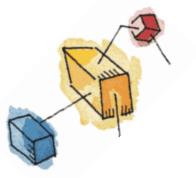
```
void writer ()
   while (true) {
     semWait (y);
          writecount++;
          if (writecount == 1) semWait (rsem);
     semSignal (y);
     semWait (wsem);
     WRITEUNIT();
     semSignal (wsem);
     semWait (y);
          writecount--;
          if (writecount == 0) semSignal (rsem);
     semSignal (y);
void main()
   readcount = writecount = 0;
   parbegin (reader, writer);
```

## State of the Process Queues

Readers only in the system	•wsem set •no queues
Writers only in the system	•wsem and rsem set •writers queue on wsem
Both readers and writers with read first	•wsem set by reader •rsem set by writer •all writers queue on wsem •one reader queues on rsem •other readers queue on z
Both readers and writers with write first	•wsem set by writer •rsem set by writer •writers queue on wsem •one reader queues on rsem •other readers queue on z







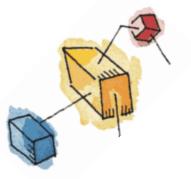
#### Summary

- Operating system themes are:
  - Multiprogramming, multiprocessing, distributed processing
  - Fundamental to these themes is concurrency
    - Issues of conflict resolution and cooperation arise
    - Effective solution: mutual exclusion, no deadlock and starvation

#### Mutual exclusion

- Condition in which there is a set of concurrent processes, only one of which is able to access a given resource or perform a given function at any time
- One approach involves the use of special purpose machine instructions – may not be efficient as it uses busy waiting





#### Summary

#### Semaphores

 Used for signalling among processes and can be readily used to enforce a mutual exclusion discipline

#### Monitors

 A programming-language construct that provides equivalent functionality to that of semaphores and that is sometimes easier to control

#### Messages

- Useful for the enforcement of synchronization discipline
- Exchange information



