

The slide features a decorative design with thin blue lines. A vertical line runs down the left side, and a horizontal line runs across the top. These lines intersect at the top-left corner, where a small blue circle is positioned. Another horizontal line runs across the bottom, and a vertical line runs down the right side. These lines intersect at the bottom-right corner, where another small blue circle is positioned. The title text is centered between these lines.

Programming Shared Address Space Platforms (Pthreads)

Thread Basics

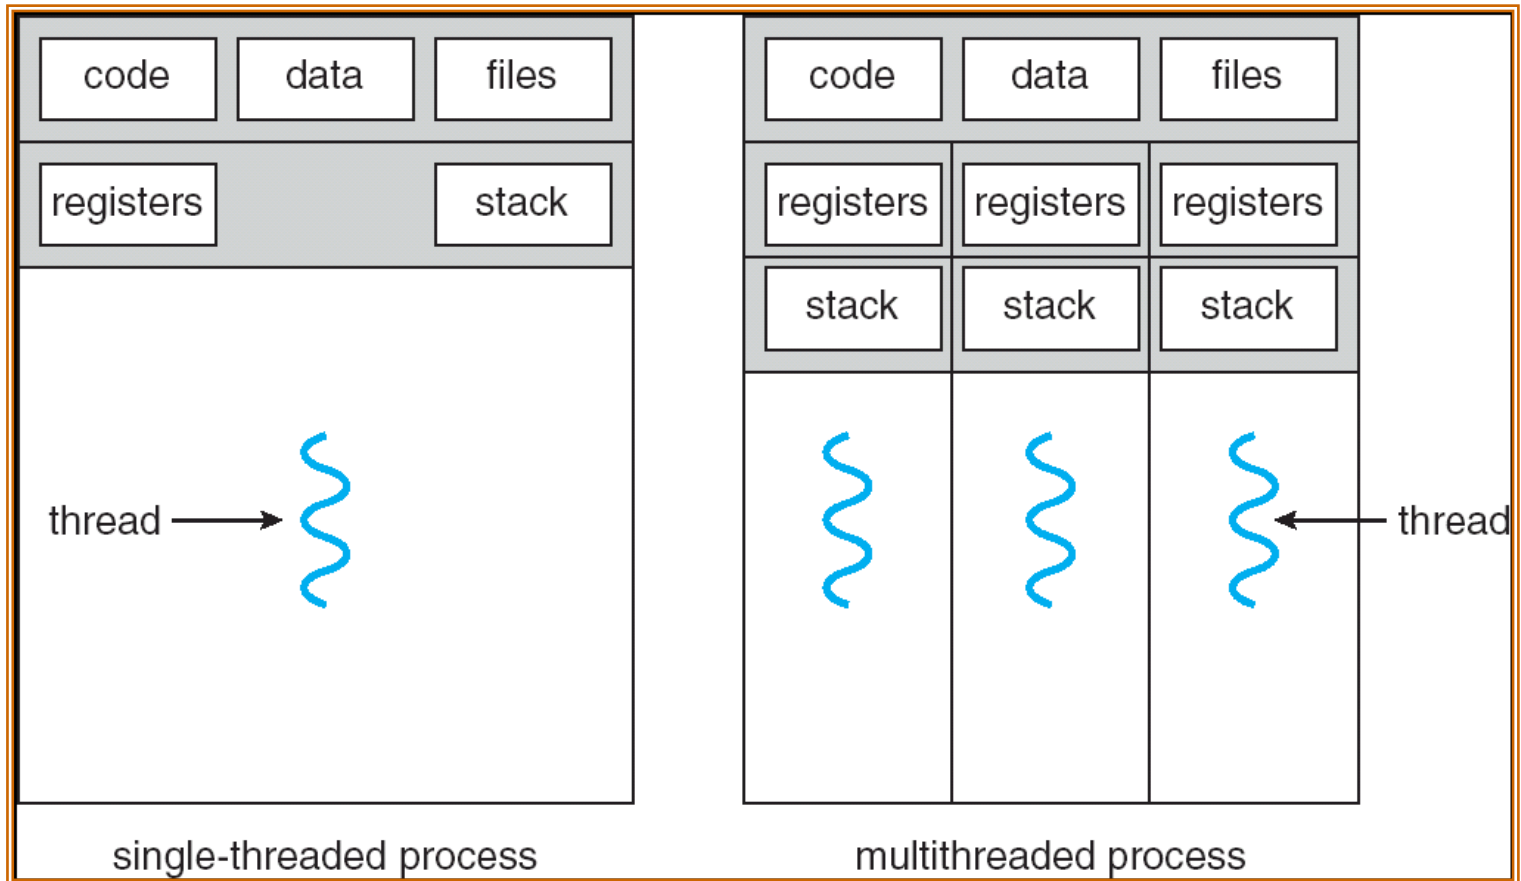


The logical machine model of a thread-based parallel programming paradigm.

Thread Basics

- ◆ All memory in the logical machine model of a thread is globally accessible to every other thread.
- ◆ The stack corresponding to the function call is generally treated as being local to the thread for liveness reasons.

Thread Basics



Thread Basics

- ❖ Cannot assume any order of execution.
- ❖ Any such order must be **explicitly** established using synchronization mechanisms.
- ❖ Issues on deadlock, starvation and performance.

Pthreads API

- ◆ The Pthreads API is defined in the ANSI/IEEE POSIX 1003.1 - 1995 standard.
- ◆ The subroutines which comprise the Pthreads API can be informally grouped into three major classes:
 - Thread management
 - Mutexes
 - Condition variables

Thread Management

- ◆ The first class of functions work directly on threads - creating, detaching, joining, etc.
- ◆ They also include functions to set/query thread attributes (joinable, scheduling etc.)

Creating Threads

- ◆ Initially, only a single, default thread. All other threads must be **explicitly created** by the programmer.
- ◆ **routine:** `pthread_create(thread, attr, start_routine, arg)`
 - **thread:** unique identifier for the new thread (*pthread_t*)
 - **attr:** attribute object used to set thread attributes. (*pthread_attr_t*) You can specify a thread attributes object, or NULL for the default values.
 - **start_routine:** C routine that the thread will execute.
 - **arg:** A single argument that may be passed to *start_routine*. It must be passed by reference. NULL may be used if no argument is to be passed.
- ◆ If successful, the `pthread_create()` function shall return zero; otherwise, an error number shall be returned to indicate the error.

Thread Attributes

- ◆ By default, a thread is created with certain attributes. Some of these attributes can be changed by the programmer via the thread attribute object.
- ◆ `pthread_attr_init(attr)` and `pthread_attr_destroy(attr)` are used to initialize/destroy the thread attribute object.
- ◆ Other routines are then used to query/set specific attributes in the thread attribute object.

Terminating Thread

- ◆ There are several ways in which a Pthread may be terminated:
 - The thread makes a call to the `pthread_exit()` subroutine.
 - The thread is cancelled by another thread via the `pthread_cancel()` routine.
 - The entire process is terminated due to a call to the `exit` subroutine.

Terminating Thread

- ◆ **routine:** `pthread_exit(status)`
- ◆ used to explicitly exit a thread.
- ◆ The programmer may optionally specify a termination status, which is stored as a void pointer for any thread that may join the calling thread.
- ◆ Cleanup: the `pthread_exit()` routine does not close files; any files opened inside the thread will remain open after the thread is terminated.

Example

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid) {
    int *tid;
    tid = (int *) threadid;
    printf("Hello World! It's me, thread #%-d!\n", *tid);
    pthread_exit(NULL);
}
```

Example

```
int main(int argc, char *argv[]) {  
    pthread_t threads[NUM_THREADS];  
    int rc, t, tids[NUM_THREADS];  
    for(t=0;t<NUM_THREADS;t++){  
        printf("In main: creating thread %d\n", t);  
        tids[t] = t;  
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *) &tids[t]);  
        if (rc){  
            printf("ERROR; return code from pthread_create() is %d\n", rc);  
            exit(-1);  
        }  
    }  
  
    pthread_exit(NULL);  
}
```

Passing Arguments to Threads

- ◆ The `pthread_create()` routine permits the programmer to pass one argument to the thread start routine.
- ◆ For cases where multiple arguments must be passed, this limitation is easily overcome by creating a structure which contains all of the arguments, and then passing a pointer to that structure in the `pthread_create()` routine.
- ◆ All arguments must be passed by reference and cast to `(void *)`.

Passing Arguments to Threads

```
struct two_args {  
    int arg1;  
    int arg2;  
};  
void *needs_2_args(void *ap) {  
    struct two_args *argp;  
    int a1, a2;  
  
    ...  
    argp = (struct two_args *) ap;  
    a1 = argp->arg1;  
    a2 = argp->arg2;  
  
    ...  
    free (argp);  
    pthread_exit(NULL);  
}
```

Passing Arguments to Threads

```
int main(int argc, char *argv[]) {  
    pthread_t t;  
    struct two_args *ap;  
    int rc;  
  
    ...  
    ap = (struct two_args *) malloc (sizeof (struct two_args));  
    ap->arg1 = 1;  
    ap->arg2 = 2;  
    rc = pthread_create(&t, NULL, needs_2_args, (void *) ap);  
    ...  
    pthread_exit(NULL);  
}
```


Joining & Detaching Threads

◆ Routines:

`pthread_join (threadid,status)`

`pthread_detach (threadid,status)`

`pthread_attr_setdetachstate (attr,detachstate)`

`pthread_attr_getdetachstate (attr,detachstate)`

◆ "Joining" is one way to accomplish synchronization between threads.

◆ The `pthread_join()` subroutine blocks the calling thread until the specified threadid thread terminates.

◆ The programmer is able to obtain the target thread's termination return status if it was specified in the target thread's call to `pthread_exit()`.

Joining & Detaching Threads

- ◆ When a thread is created, one of its attributes defines whether it is joinable or detached.
- ◆ Only threads that are created as joinable can be joined. If a thread is created as detached, it can never be joined.
- ◆ To explicitly create a thread as joinable or detached, the attr argument in the `pthread_create()` routine is used:
 - Declare a pthread attribute variable of the `pthread_attr_t` data type
 - Initialize the attribute variable with `pthread_attr_init()`
 - Set the attribute detached status with `pthread_attr_setdetachstate()`
 - When done, free library resources used by the attribute with `pthread_attr_destroy()`

Example

```
void *BusyWork(void *null) {  
    ...  
    pthread_exit((void *) 0);  
}
```

Example

```
int main (int argc, char *argv[]) {
    pthread_attr_t attr;
    int rc, t;
    void *status;
    /* Initialize and set thread detached attribute */
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);
    ...
    /* Free attribute and wait for the other threads */
    pthread_attr_destroy(&attr);
    for(t=0; t<NUM_THREADS; t++) {
        rc = pthread_join(thread[t], &status);
        ...
        printf("Completed join with thread %d status= %ld\n",t,
(long)status);
    }
    pthread_exit(NULL);
}
```

Synchronization Issues

- ◆ When multiple threads attempt to manipulate the same data item, the results can often be incoherent if proper care is not taken to synchronize them.
- ◆ Consider:

```
/* each thread tries to update variable best_cost as follows */  
if (my_cost < best_cost)  
    best_cost = my_cost;
```
- ◆ Assume that there are two threads, the initial value of `best_cost` is 100, and the values of `my_cost` are 50 and 75 at threads `t1` and `t2`.
- ◆ Depending on the schedule of the threads, the value of `best_cost` could be 50 or 75 - A **race condition** problem!

Race Condition

- ◆ A race condition occurs when
 - Multiple processes or threads read and write shared data items
 - They do so in a way where the final result depends on the order of execution of the processes.
- ◆ The output depends on who finishes the race last.
- ◆ Cannot assume any order of execution.
- ◆ Any such order must be explicitly established using synchronization mechanisms.

Mutex

- ◆ The second class of functions deal with synchronization, called a "mutex", which is an abbreviation for "mutual exclusion".
- ◆ Mutex functions provide for creating, destroying, locking and unlocking mutexes.
- ◆ They are also supplemented by mutex attribute functions that set or modify attributes associated with mutexes.
- ◆ When a process, or thread executes code that manipulates shared data, we say it is in its **Critical Section**.
- ◆ A general structure for critical section:

```
...  
lock mutex  
    critical section  
unlock mutex  
noncritical section
```

Creating & Destroying Mutex

- ◆ Mutex must be declared with type `pthread_mutex_t`, and must be initialized before they can be used.
- ◆ There are two ways to initialize a mutex variable:
 - Statically, when it is declared. For example:
`pthread_mutex_t mymutex = PTHREAD_MUTEX_INITIALIZER;`
 - Dynamically, with the `pthread_mutex_init()` routine. This method permits setting mutex object attributes, *attr* (which may be specified as `NULL` to accept defaults).
- ◆ **Routines:**
 - `pthread_mutex_init (mutex, attr)`
 - `pthread_mutex_destroy (mutex)`
 - `pthread_mutexattr_init (attr)`
 - `pthread_mutexattr_destroy (attr)`
- ◆ The mutex is initially unlocked.

Locking & Unlocking Mutex

◆ Routines:

pthread_mutex_lock (*mutex*)

pthread_mutex_unlock (*mutex*)

pthread_mutex_trylock (*mutex*)

- ◆ pthread_mutex_lock() is used by a thread to acquire a lock on the specified *mutex* variable.
- ◆ pthread_mutex_unlock() will unlock a mutex if called by the owning thread. An error will be returned if:
 - the mutex was already unlocked
 - the mutex is owned by another thread
- ◆ pthread_mutex_trylock() will attempt to lock a mutex. However, if the mutex is already locked, the routine will return immediately with an "EBUSY" error code. This routine may be useful in preventing deadlock conditions.

Example 1

◆ We can now write our previously incorrect code segment as:

◆

```
....
pthread_mutex_t minimum_value_lock;
int best_cost; //global variable!
main() {
    ....
    pthread_mutex_init(&minimum_value_lock, NULL);
    ....
}
void *find_min(void *list_ptr) {
    ....
    pthread_mutex_lock(&minimum_value_lock);
    if (my_cost < best_cost)
        best_cost = my_cost;
    pthread_mutex_unlock(&minimum_value_lock);
}
```

Example 2

- ◆ The **producer-consumer scenario** imposes the following constraints:
 - The producer thread must not overwrite the shared buffer when the previous task has not been picked up by a consumer thread.
 - The consumer threads must not pick up tasks until there is something present in the shared data structure.
 - Individual consumer threads should pick up tasks one at a time.

Example 2

```
pthread_mutex_t task_queue_lock;  
int task_available;  
...  
main() {  
    ....  
    task_available = 0;  
    pthread_mutex_init(&task_queue_lock, NULL);  
    ....  
}
```

Example 2

```
void *producer(void *producer_thread_data) {  
    ....  
    while (!done()) {  
        inserted = 0;  
        create_task(&my_task);  
        while (inserted == 0) {  
            pthread_mutex_lock(&task_queue_lock);  
            if (task_available == 0) {  
                insert_into_queue(my_task);  
                task_available = 1;  
                inserted = 1;  
            }  
            pthread_mutex_unlock(&task_queue_lock);  
        }  
    }  
}
```

Example 2

```
void *consumer(void *consumer_thread_data) {  
    ...  
    while (!done()) {  
        extracted = 0;  
        while (extracted == 0) {  
            pthread_mutex_lock(&task_queue_lock);  
            if (task_available == 1) {  
                extract_from_queue(&my_task);  
                task_available = 0;  
                extracted = 1;  
            }  
            pthread_mutex_unlock(&task_queue_lock);  
        }  
        process_task(my_task);  
    }  
}
```

Overheads of Locking

- ◆ Locks represent serialization points since critical sections must be executed by threads one after another.
- ◆ Encapsulating large segments of the program within locks can lead to significant performance degradation.
- ◆ It is often possible to reduce the idling overhead associated with locks using `pthread_mutex_trylock`.

Alleviating Locking Overhead

```
/* using pthread_mutex_trylock routine */
...
Pthread_mutex_t tryLock_lock = PTHREAD_MUTEX_INITIALIZER;
...
lock_status=pthread_mutex_trylock(&tryLock_lock);
if (lock_status == EBUSY) {
    /* do something else */
    ...
}
else {
    /* do one thing */
    ...
    pthread_mutex_unlock(&tryLock_lock);
}
...
```


Condition Variables

- ◆ The third class of functions address communications between threads that share a mutex.
- ◆ A condition variable allows a thread to block itself until specified data reaches a predefined state.
- ◆ A condition variable indicates an event and has no value.
 - One cannot store a value into nor retrieve a value from a condition variable.
 - If a thread must wait for an event to occur, that thread **waits** on the corresponding condition variable.
 - A condition variable has a queue for those threads that are waiting the corresponding event to occur to wait on.
 - If another thread causes the event to occur, that thread simply **signals** the corresponding condition variable.

Condition Variables

- ◆ This class includes functions to create, destroy, wait and signal based upon specified variable values.
- ◆ Functions to set/query condition variable attributes are also included.
- ◆ A condition variable is **always used** in conjunction **with a mutex lock**.

Creating & Destroying Condition

Variables

- ◆ Condition variables must be declared with type `pthread_cond_t`, and must be initialized before they can be used.
- ◆ There are two ways to initialize a condition variable:
 - Statically, when it is declared. For example:
`pthread_cond_t myconvar = PTHREAD_COND_INITIALIZER;`
 - Dynamically, with the `pthread_cond_init()` routine. This method permits setting condition variable object attributes, *attr* (which may be specified as `NULL` to accept defaults).

◆ Routines:

`pthread_cond_init (condition, attr)`

`pthread_cond_destroy (condition)`

`pthread_condattr_init (attr)`

`pthread_condattr_destroy (attr)`

Waiting and Signaling on Condition

Variables

◆ Routines:

pthread_cond_wait (*condition*, *mutex*)

pthread_cond_signal (*condition*)

pthread_cond_broadcast (*condition*)

- ◆ pthread_cond_wait() blocks the calling thread until the specified *condition* is signalled.
- ◆ This routine should be called while *mutex* is locked, and it will automatically release the *mutex* while it waits.
- ◆ After signal is received and thread is awakened, *mutex* will be automatically locked for use by the thread.
- ◆ The programmer is then responsible for unlocking *mutex* when the thread is finished with it.

Waiting and Signaling on Condition

Variables

- ◆ `pthread_cond_signal()` is used to signal (or wake up) another thread which is waiting on the condition variable.
 - It should be called after *mutex* is locked, and
 - must unlock *mutex* in order for `pthread_cond_wait()` routine to complete.
- ◆ The `pthread_cond_broadcast()` routine unlocks all of the threads blocked on the condition variable.

Waiting and Signaling on Condition

Variables

◆ Proper locking and unlocking of the associated *mutex* variable is essential when using these routines. For example:

- Failing to lock the *mutex* before calling `pthread_cond_wait()` may cause it NOT to block.
- Failing to unlock the *mutex* after calling `pthread_cond_signal()` may not allow a matching `pthread_cond_wait()` routine to complete (it will remain blocked).

Producer-Consumer Using Condition Variables

```
pthread_cond_t cond_queue_empty, cond_queue_full;
pthread_mutex_t task_queue_cond_lock;
int task_available;
/* other data structures here */
...
main() {
    /* declarations and initializations */
    task_available = 0;
    pthread_cond_init(&cond_queue_empty, NULL);
    pthread_cond_init(&cond_queue_full, NULL);
    pthread_mutex_init(&task_queue_cond_lock, NULL);
    /* create and join producer and consumer threads */
    ...
}
```

Producer-Consumer Using Condition Variables

```
void *producer(void *producer_thread_data) {  
    while (!done()) {  
        create_task();  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (task_available == 1)  
            pthread_cond_wait(&cond_queue_empty,  
                             &task_queue_cond_lock);  
        insert_into_queue();  
        task_available = 1;  
        pthread_cond_signal(&cond_queue_full);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
    }  
}
```


Producer-Consumer Using Condition Variables

```
void *consumer(void *consumer_thread_data) {  
    while (!done()) {  
        pthread_mutex_lock(&task_queue_cond_lock);  
        while (task_available == 0)  
            pthread_cond_wait(&cond_queue_full,  
                             &task_queue_cond_lock);  
        my_task = extract_from_queue();  
        task_available = 0;  
        pthread_cond_signal(&cond_queue_empty);  
        pthread_mutex_unlock(&task_queue_cond_lock);  
        process_task(my_task);  
    }  
}
```